

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv("F:\\NEW_DATSET\\CarPrice_Assignment.csv")
df.shape
```

Out[2]: (205, 26)

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                 205 non-null    int64
1   symboling              205 non-null    int64
2   CarName                205 non-null    object
3   fueltype               205 non-null    object
4   aspiration              205 non-null    object
5   doornumber              205 non-null    object
6   carbody                205 non-null    object
7   drivewheel             205 non-null    object
8   enginelocation          205 non-null    object
9   wheelbase              205 non-null    float64
10  carlength              205 non-null    float64
11  carwidth                205 non-null    float64
12  carheight              205 non-null    float64
13  curbweight              205 non-null    int64
14  enginetype              205 non-null    object
15  cylindernumber          205 non-null    object
16  enginesize              205 non-null    int64
17  fuelsystem              205 non-null    object
18  boreratio               205 non-null    float64
19  stroke                  205 non-null    float64
20  compressionratio        205 non-null    float64
21  horsepower              205 non-null    int64
22  peakrpm                 205 non-null    int64
23  citympg                 205 non-null    int64
24  highwaympg              205 non-null    int64
25  price                   205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

```
In [4]: df.describe()
```

```
Out[4]:
```


	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionr.
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	126.907317	3.329756	3.255415	10.142
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	41.642693	0.270844	0.313597	3.972
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.000000	2.540000	2.070000	7.000
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97.000000	3.150000	3.110000	8.600
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.000000	3.310000	3.290000	9.000
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141.000000	3.580000	3.410000	9.400
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326.000000	3.940000	4.170000	23.000

```
In [5]: df.fuelsystem.unique()
```

```
Out[5]: array(['mpfi', '2bbl', 'mfi', '1bbl', 'spfi', '4bbl', 'idi', 'spdi'],
dtype=object)
```

```
In [6]: plt.scatter(x=df.price,y=df.enginesize)
```

Out[6]: <matplotlib.collections.PathCollection at 0x24982789e20>




```
In [7]: q1,q2=np.percentile(df['price'],25),np.percentile(df['price'],75)
cutt_off=(q2-q1)*1.5
lower,upper=q1,q2+cutt_off
print(lower,upper)
df=df[(df['price']<upper)&(df['price']>lower)]

1 29575.5
```


```
In [8]: sns.boxplot(x='price',data=df)
```

Out[8]: <AxesSubplot:xlabel='price'>



```
In [9]: df.price.plot()
```

Out[9]: <AxesSubplot:>



In [ ]:

```
In [10]: df.isnull().sum()
```

```
Out[10]: car_ID          0
symboling          0
CarName           0
fueltype          0
aspiration        0
doornumber        0
carbody           0
drivewheel        0
enginelocation    0
wheelbase         0
carlength         0
carwidth          0
carheight         0
curbweight        0
engine           0
cylindernumber    0
enginesize        0
fuelsystem        0
boreratio         0
stroke            0
compressionratio  0
horsepower        0
peakrpm           0
citympg           0
highwaympg        0
price             0
dtype: int64
```

```
In [11]: df=df.drop(['CarName'],axis=1)
```

```
In [12]: df.fueltype.value_counts()
```

```
Out[12]: gas          171
diesel          19
Name: fueltype, dtype: int64
```

```
In [13]: fueltype=pd.get_dummies(df['fueltype'],drop_first=True)
fueltype.head()
```

```
Out[13]:   gas
0      1
1      1
2      1
3      1
4      1
```

```
In [14]: carbod=pd.get_dummies(df['carbody'],drop_first=True)
carbod.tail()
```

```
Out[14]:   hardtop  hatchback  sedan  wagon
200      0         0         1         0
201      0         0         1         0
202      0         0         1         0
203      0         0         1         0
204      0         0         1         0
```

```
In [15]: engineloc=pd.get_dummies(df['enginelocation'],drop_first=True)
drivewheels=pd.get_dummies(df['drivewheel'],drop_first=True)
drivewheels.head()
```

```
Out[15]:   fwd  rwd
0      0   1
```

1	0	1
2	0	1
3	1	0
4	0	0

```
In [16]: gassystem=pd.get_dummies(df['fuelsystem'],drop_first=True)
gassystem.head()
```

```
Out[16]:
```

	2bbl	4bbl	idi	mfi	mpfi	spdi	spfi
0	0	0	0	0	1	0	0
1	0	0	0	0	1	0	0
2	0	0	0	0	1	0	0
3	0	0	0	0	1	0	0
4	0	0	0	0	1	0	0

```
In [17]: doorno=pd.get_dummies(df['doornumber'],drop_first=True)
doorno.head()
```

```
Out[17]:
```

	two
0	1
1	1
2	1
3	0
4	0

```
In [18]: df.head()
```

```
Out[18]:
```

	car_ID	symboling	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	...	enginesize	fuelsystem
0	1	3	gas	std	two	convertible	rwd	front	88.6	168.8	...	130	mpfi
1	2	3	gas	std	two	convertible	rwd	front	88.6	168.8	...	130	mpfi
2	3	1	gas	std	two	hatchback	rwd	front	94.5	171.2	...	152	mpfi
3	4	2	gas	std	four	sedan	fwd	front	99.8	176.6	...	109	mpfi
4	5	2	gas	std	four	sedan	4wd	front	99.4	176.6	...	136	mpfi

5 rows × 25 columns

```
In [19]: df=df.drop(['aspiration','carbody','enginelocation','drivewheel','fuelsystem','doornumber','fueltype'],axis=1)
df.head()
```

```
Out[19]:
```

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginetype	cylindernumber	enginesize	boreratio	stroke	compre
0	1	3	88.6	168.8	64.1	48.8	2548	dohc	four	130	3.47	2.68	
1	2	3	88.6	168.8	64.1	48.8	2548	dohc	four	130	3.47	2.68	
2	3	1	94.5	171.2	65.5	52.4	2823	ohcv	six	152	2.68	3.47	
3	4	2	99.8	176.6	66.2	54.3	2337	ohc	four	109	3.19	3.40	
4	5	2	99.4	176.6	66.4	54.3	2824	ohc	five	136	3.19	3.40	

```
In [20]: cylinderno=pd.get_dummies(df['cylindernumber'],drop_first=True)
cylinderno.head()
```

```
Out[20]:
```

	four	six	three	two
0	1	0	0	0
1	1	0	0	0
2	0	1	0	0

3	1	0	0	0
4	0	0	0	0

```
In [21]: df=pd.concat([cylinderno,df],axis=1)
df.isnull().sum()
```


```
Out[21]: four          0
six          0
three        0
two          0
car_ID       0
symboling    0
wheelbase    0
carlength    0
carwidth     0
carheight    0
curbweight   0
enginetype   0
cylindernumber 0
enginesize   0
boreratio    0
stroke       0
compressionratio 0
horsepower   0
peakrpm      0
citympg      0
highwaympg   0
price        0
dtype: int64
```

```
In [22]: df=df.drop(['cylindernumber'],axis=1)
df.head()
```

```
Out[22]:
```

	four	six	three	two	car_ID	symboling	wheelbase	carlength	carwidth	carheight	...	enginetype	enginesize	boreratio	stroke	compressi
0	1	0	0	0	1	3	88.6	168.8	64.1	48.8	...	dohc	130	3.47	2.68	
1	1	0	0	0	2	3	88.6	168.8	64.1	48.8	...	dohc	130	3.47	2.68	
2	0	1	0	0	3	1	94.5	171.2	65.5	52.4	...	ohcv	152	2.68	3.47	
3	1	0	0	0	4	2	99.8	176.6	66.2	54.3	...	ohc	109	3.19	3.40	
4	0	0	0	0	5	2	99.4	176.6	66.4	54.3	...	ohc	136	3.19	3.40	

5 rows × 21 columns



```
In [23]: enginetype=pd.get_dummies(df['enginetype'],drop_first=True)
enginetype.head()
```

```
Out[23]:
```

	l	ohc	ohcf	ohcv	rotor
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	1	0
3	0	1	0	0	0
4	0	1	0	0	0

```
In [24]: df=pd.concat([enginetype,df],axis=1)
df.isnull().sum()
```

```
Out[24]: l          0
ohc          0
ohcf         0
ohcv         0
rotor        0
four         0
six          0
three        0
two          0
```

```

car_ID      0
symboling   0
wheelbase   0
carlength   0
carwidth    0
carheight   0
curbweight  0
enginetype  0
enginesize  0
bore ratio  0
stroke      0
compressionratio  0
horsepower  0
peakrpm     0
citympg     0
highwaympg  0
price       0
dtype: int64

```

```
In [25]: df=df.drop(['enginetype'],axis=1)
```

```
In [26]: df.head()
```

```
Out[26]:
```

	l	ohc	ohcf	ohcv	rotor	four	six	three	two	car_ID	...	curbweight	enginesize	bore ratio	stroke	compressionratio	horsepower	peakrpm
0	0	0	0	0	0	1	0	0	0	1	...	2548	130	3.47	2.68	9.0	111	50
1	0	0	0	0	0	1	0	0	0	2	...	2548	130	3.47	2.68	9.0	111	50
2	0	0	0	1	0	0	1	0	0	3	...	2823	152	2.68	3.47	9.0	154	50
3	0	1	0	0	0	1	0	0	0	4	...	2337	109	3.19	3.40	10.0	102	55
4	0	1	0	0	0	0	0	0	0	5	...	2824	136	3.19	3.40	8.0	115	55

5 rows × 25 columns

```
In [27]: df.describe()
```


```
Out[27]:
```

	l	ohc	ohcf	ohcv	rotor	four	six	three	two	car_ID	...	curbweight	enginesize	bore ratio	stroke	compressionratio	horsepower	peakrpm
count	190.000000	190.000000	190.000000	190.000000	190.000000	190.000000	190.000000	190.000000	190.000000	190.000000	...	190.000000	190.000000	190.000000	190.000000	190.000000	190.000000	190.000000
mean	0.063158	0.757895	0.063158	0.042105	0.021053	0.836842	0.084211	0.005263	0.021053	105.463158	...	2480.031579	130.000000	3.470000	2.680000	9.000000	111.000000	50.000000
std	0.243889	0.429489	0.243889	0.201360	0.143939	0.370486	0.278437	0.072548	0.143939	59.925052	...	445.890591	13.000000	0.310000	0.540000	0.000000	15.000000	5.000000
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.000000	...	1488.000000	100.000000	2.680000	2.680000	8.000000	102.000000	50.000000
25%	0.000000	1.000000	0.000000	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000	54.250000	...	2128.000000	100.000000	3.190000	3.400000	9.000000	111.000000	50.000000
50%	0.000000	1.000000	0.000000	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000	106.500000	...	2395.000000	100.000000	3.190000	3.400000	9.000000	111.000000	50.000000
75%	0.000000	1.000000	0.000000	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000	157.750000	...	2823.750000	100.000000	3.190000	3.400000	9.000000	111.000000	50.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	205.000000	...	3750.000000	152.000000	3.470000	3.400000	10.000000	154.000000	55.000000

8 rows × 25 columns

```
In [28]: sns.heatmap(df[["symboling","curbweight","enginesize","bore ratio","stroke","compressionratio","horsepower","peakrpm","citympg","highwaympg","price"]],
plt.figure(figsize=(100,100))
```

```
Out[28]: <Figure size 7200x7200 with 0 Axes>
```



symbc  
 curbwe  
 engine  
 bore  
 st  
 compression  
 horsep  
 peak  
 city  
 highway  
 f

<Figure size 7200x7200 with 0 Axes>

```
In [29]: df=df.drop(['symboling','peakrpm','highwaympg','citympg','stroke'],axis=1)
```

```
In [30]: x=df.drop(['price'],axis=1)
y=df['price']
```

```
In [31]: x.shape
```

Out[31]: (190, 19)

```
In [32]: from sklearn.model_selection import train_test_split
x_tr,x_te,y_tr,y_te=train_test_split(x,y,test_size=0.15,random_state=200)
```


```
In [33]: from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_tr=sc.fit_transform(x_tr)
x_te=sc.transform(x_te)
```

```
In [34]: from sklearn.linear_model import LinearRegression
model=LinearRegression()
model.fit(x_tr,y_tr)
```

Out[34]: LinearRegression()

```
In [35]: plt.scatter(df.price,df.enginesize)
```

Out[35]: <matplotlib.collections.PathCollection at 0x249837fa8b0>



```
In [36]: y_pred=model.predict(x_te)
```

```
In [37]: y_pred
```

```
Out[37]: array([ 9791.77546556, 11332.03184687,  7527.94188211, 14722.92462009,
        17621.90830733, 13262.03776563, 19064.81182464,  6215.42094522,
        16000.47771019, 10494.94351621,  6923.48284459,  9883.86780493,
        6908.27158035, 11304.72153732, 15779.95592913,  6636.45566505,
        11355.1996256 , 15441.07000946, 17259.21935798,  5909.5160443 ,
        6697.25389125,  7776.36578613, 13314.10699518,  7672.45992822,
        6084.68701374, 13526.7635203 , 22695.28487201,  6887.84267262,
        10351.75369187])
```

```
In [38]: model.score(x_te,y_te)
```

Out[38]: 0.856182405140828

```
In [39]: print(model.intercept_)  
print(model.coef_)
```

```
11617.796068322985  
[ -767.1467782    335.86913948 -445.18141048 -643.49090382  
 -259.53509202 -985.6587882    868.61653579  265.19689594  
 -259.53509202 -824.58213014   543.60308068 -504.70448004  
 1247.55206498 -190.06077879  3106.44591898 -1826.26646283  
 1230.53681948   611.98160241  1106.01633542]
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js